

Detailed Report on Implementing a Neural Network for MNIST Classification

Luna Schätzle

December 12, 2024

Task Description

In Task 5, we implemented a neural network to classify the MNIST database. The MNIST database consists of 60,000 training and 10,000 testing images. Each image is 28×28 pixels, representing a handwritten digit between 0 and 9. The objective was to build a neural network capable of accurately classifying these digits.

The task required implementing a complete workflow, including data pre-processing, designing the neural network architecture, training the model, evaluating its performance, and utilizing the trained model for predictions.

Data Loading

The first step was to load the MNIST dataset for use in the model. The function `load_data()` was implemented to load the dataset and return the training and test sets.

Code Implementation

```
1 import numpy as np
2 import tensorflow as tf
3 from tensorflow.keras.datasets import mnist
4
5 # Load MNIST dataset
6 (x_train, y_train), (x_test, y_test) = mnist.load_data()
7
```

```
8 # Print dataset dimensions
9 print(f"x_train shape: {x_train.shape}")
10 print(f"y_train shape: {y_train.shape}")
11 print(f"x_test shape: {x_test.shape}")
12 print(f"y_test shape: {y_test.shape}")
```

Dataset Analysis

The MNIST dataset consists of grayscale images, where pixel values range from 0 to 255. To simplify processing for the neural network, pixel values were normalized to the range $[0, 1]$ by dividing by 255. This normalization ensures the neural network can process data efficiently without encountering issues related to large input values.

Visualization

An example image from the training dataset is displayed below, alongside sample images from each digit category:




Figure 1: Example Image from Training Dataset

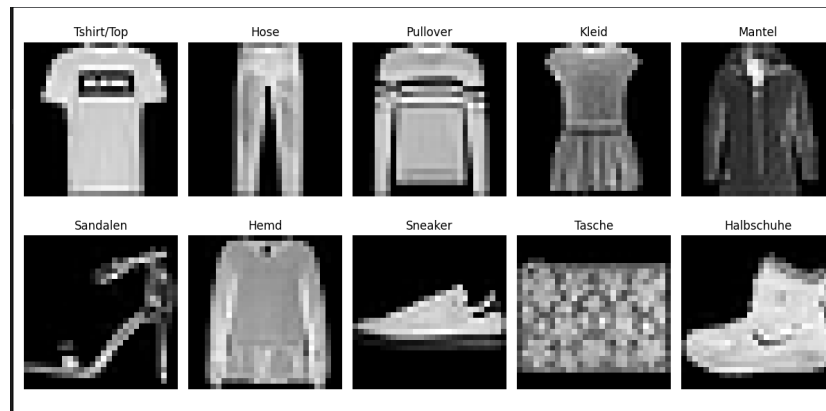


Figure 2: Sample Images from All Categories

Model Architecture

The neural network was implemented using Keras. The architecture includes:

- An input layer for 28×28 flattened pixels.
- Six dense (fully connected) layers with ReLU activation.
- A final output layer with softmax activation for 10 classes.

This architecture was chosen to balance simplicity and performance. By stacking multiple layers with decreasing sizes, the model learns progressively abstract features, which aids in accurate classification.

Code Implementation

```

1 from tensorflow.keras import models, layers
2
3 model = models.Sequential([
4     layers.Input(shape=(28 * 28,)),
5     layers.Dense(512, activation='relu'),
6     layers.Dense(256, activation='relu'),
7     layers.Dense(128, activation='relu'),
8     layers.Dense(24, activation='relu'),
9     layers.Dense(10, activation='relu'),
10    layers.Dense(10, activation='softmax')
11 ])

```

Model Summary

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 512)	401,920
dense_1 (Dense)	(None, 256)	131,328
dense_2 (Dense)	(None, 128)	32,896
dense_3 (Dense)	(None, 24)	3,096
dense_4 (Dense)	(None, 10)	250
dense_5 (Dense)	(None, 10)	110

Total params: 569,600 (2.17 MB)

Trainable params: 569,600 (2.17 MB)

Non-trainable params: 0 (0.00 B)

Training Process

The model was trained over 40 epochs using the Adam optimizer and sparse categorical cross-entropy loss. Training and validation accuracy and loss were tracked.

Training Details

Training involved the following steps:

1. Compiling the model with the Adam optimizer and sparse categorical cross-entropy loss.
2. Splitting data into training and validation sets to monitor performance.
3. Running the training process for 40 epochs with a batch size of 32.

Training Output

The training process demonstrated gradual improvements in accuracy and reductions in loss over epochs, as shown in [Figure 4](#).




Figure 3: Training Accuracy and Loss over Epochs

Evaluation

After training, the model was evaluated on the test set. The results were:

- Test Loss: 0.5342
- Test Accuracy: 88.74%

Performance Insights

The model achieved high accuracy, indicating effective learning. However, further improvements could involve:

- Adding dropout layers to reduce overfitting.
- Using convolutional layers to capture spatial relationships.
- Tuning hyperparameters such as learning rate and batch size.

Model Usage

The trained model was saved as `fashion_mnist_model.h5` and can be loaded for predictions. A simple prediction function was implemented to classify images:

```

1 def predict_image(img_array):
2     prediction = model.predict(img_array)
3     predicted_class = np.argmax(prediction, axis=1)[0]
4     confidence = np.max(prediction) * 100
5     return predicted_class, confidence

```

Image Preprocessing

Images must be preprocessed before being passed into the model. Preprocessing includes:

- Converting the image to grayscale.
- Resizing the image to 28×28 pixels.
- Normalizing pixel values to the range $[0, 1]$.
- Flattening the image into a single vector.

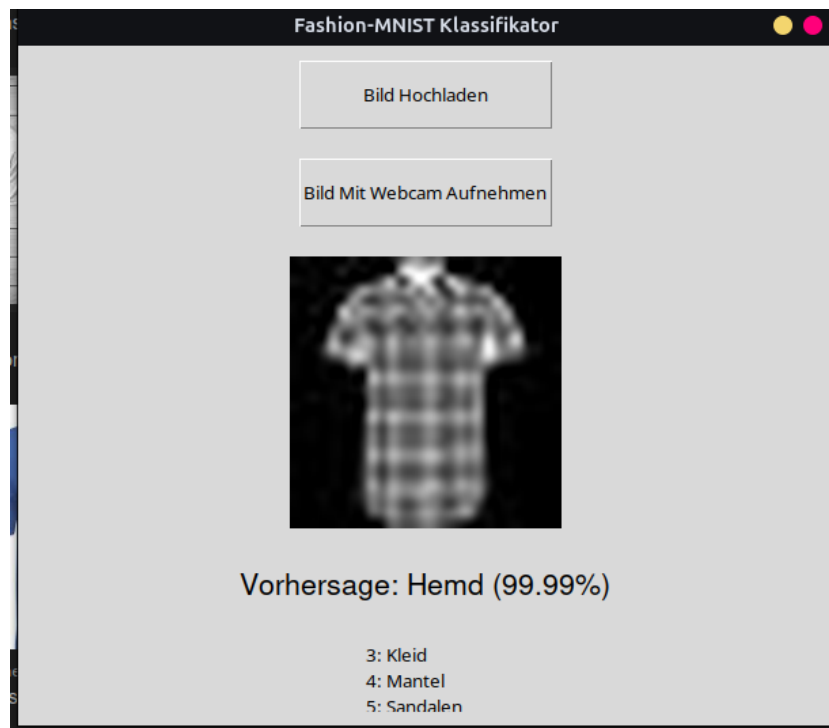


Figure 4: Implementation of Image Preprocessing and simple User Interface

Conclusion

This project successfully implemented a neural network for MNIST digit classification. With a test accuracy of 88.74%, the model demonstrates reasonable performance. Future improvements could involve hyperparameter tuning, experimenting with convolutional neural networks (CNNs), and increasing the dataset size for enhanced generalization.